

Long Homework 2

Bhavika Sewpal - 300089940

4/10/2021

Question 1 - Empirical Study: Cross-Validation Method 1

```
setwd("C:/Users/mahim/Desktop/winter2021/MAT3373/long_hw2")
mnist <- read.csv("mnist.csv")
label <- as.factor(mnist$label)

normalize <- function(x){
  if (max(x) - min(x) == 0){
    return (x)
  }else{
    return ((x-min(x))/(max(x) - min(x)))
  }
}

#normalizing the features
mnist <- as.data.frame(lapply(mnist[,2:ncol(mnist)],normalize))
mnist <- cbind(label,mnist)

set.seed(100)
index1 <- sample(c(1:10000),3333)
fold1 <- mnist[index1,]

mnist2 <- mnist[-index1,]
set.seed(100)
index2 <- sample(c(1:6667),3333)
fold2 <- mnist2[index2,]

fold3 <- mnist2[-index2,]

#Train on fold1 and fold2, test on fold3
train1 <- rbind(fold1,fold2)
train1_features <- train1[,-1]
train1_target <- train1[,1]

test1 <- fold3
test1_features <- test1[,-1]
test1_target <- test1[,1]
```

```
library(class)
knn1 <- knn(train1_features, test1_features, train1_target, k=3)
table(test1_target,knn1)
```

```
##          knn1
## test1_target  0  1  2  3  4  5  6  7  8  9
##          0 316  1  0  0  0  1  4  0  0  0
##          1  0 372  1  1  0  0  1  0  0  0
##          2  5  3 338  1  1  0  4  5  2  0
##          3  0  1  0 309  0  4  0  8  0  3
##          4  0  2  1  0 305  0  1  0  0 14
##          5  2  3  0 13  0 262  2  0  1  4
##          6  2  1  0  0  0  1 319  0  1  0
##          7  0 13  2  0  0  0  0 325  0  3
##          8  2  2  4  9  1  5  2  2 296  3
##          9  4  3  1  2  9  0  0  7  2 322
```

```
error_rate1 <- mean(knn1 != test1_target)
error_rate1
```

```
## [1] 0.0509898
```

```
#Train on fold1 and fold3, test on fold2
```

```
train2 <- rbind(fold1,fold3)
train2_features <- train2[,-1]
train2_target <- train2[,1]
```

```
test2 <- fold2
test2_features <- test2[,-1]
test2_target <- test2[,1]
```

```
knn2 <- knn(train2_features, test2_features, train2_target, k=3)
table(test2_target,knn2)
```

```
##          knn2
## test2_target  0  1  2  3  4  5  6  7  8  9
##          0 323  0  0  0  0  2  1  0  0  0
##          1  0 373  2  0  0  0  0  1  0  0
##          2  6  9 330  3  0  0  1  7  1  0
##          3  0  1  1 320  0  8  0  1  4  2
##          4  0  3  0  0 307  1  1  0  1 11
##          5  0  5  0  8  1 267  2  1  1  4
##          6  4  3  0  0  1  1 308  0  1  0
##          7  0  8  0  1  1  0  0 315  0  6
##          8  1  4  2 10  3  8  2  2 301  4
##          9  1  5  0  1  5  1  0  8  1 316
```

```
error_rate2 <- mean(knn2 != test2_target)
error_rate2
```

```
## [1] 0.05190519
```

```

#Train on fold2 and fold3, test on fold1
train3 <- rbind(fold2,fold3)
train3_features <- train3[,-1]
train3_target <- train3[,1]

test3 <- fold1
test3_features <- test3[,-1]
test3_target <- test3[,1]

knn3 <- knn(train3_features, test3_features, train3_target, k=3)
table(test3_target,knn3)

```

```

##           knn3
## test3_target  0   1   2   3   4   5   6   7   8   9
##           0 328   0   0   0   0   0   2   1   1   0
##           1   0 382   1   0   0   0   1   0   0   0
##           2   2   6 291   1   1   0   1  12   1   1
##           3   0   0   2 331   0   6   1   3   5   0
##           4   0   7   0   0 313   0   4   0   0  11
##           5   2   0   0   7   0 301   4   1   1   0
##           6   1   1   0   0   0   0 314   0   0   0
##           7   1  11   2   0   2   0   1 330   0   7
##           8   1   5   4   7   2   4   2   3 280   3
##           9   0   0   2   4   5   1   0   8   2 299

```

```

error_rate3 <- mean(knn3 != test3_target)
error_rate3

```

```
## [1] 0.04920492
```

```

average <- (error_rate1 + error_rate2 + error_rate3)/ 3
average

```

```
## [1] 0.05069997
```

The error rate is 0.0507.

For hw1, the error rate was about 0.18 .

The error rate when using cross-validation is much smaller than that without cross-validation.

In general, I would have expected the contrary. Since, we are fitting the model on only 2/3 of the data, I would expect the error rate obtained empirically to be an overestimate of the test error.

Question 2 - Empirical Study: Cross-Validation Method 2

```

#function to calculate the residual standard error - sigma
rse <- function(true,pred){
  sse <- sum((pred-true)^2)
  mean_sse <- sse / (length(pred))
}

redwine <- read.csv("redwine.csv",sep=";")
set.seed(100)

```

```
test_index <- sample(c(1:1599),799)
test <- redwine[test_index,]
train <- redwine[-test_index,]
#Linear Regression
linear_model <- lm(quality~.,data=train)
summary(linear_model)
```

```
##
## Call:
## lm(formula = quality ~ ., data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.41873	-0.38160	-0.04562	0.44648	2.10450

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.1018296	30.3018271	-0.201	0.840463
fixed.acidity	0.0055089	0.0368867	0.149	0.881317
volatile.acidity	-1.1722749	0.1792308	-6.541	1.10e-10 ***
citric.acid	-0.1221569	0.2261068	-0.540	0.589169
residual.sugar	-0.0028991	0.0218773	-0.133	0.894610
chlorides	-2.3202111	0.6351444	-3.653	0.000276 ***
free.sulfur.dioxide	0.0043875	0.0030721	1.428	0.153644
total.sulfur.dioxide	-0.0033095	0.0009774	-3.386	0.000745 ***
density	11.0501078	30.9243710	0.357	0.720943
pH	-0.6007354	0.2737706	-2.194	0.028505 *
sulphates	0.9385347	0.1696135	5.533	4.28e-08 ***
alcohol	0.2882921	0.0380634	7.574	1.01e-13 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6599 on 788 degrees of freedom
## Multiple R-squared:  0.3667, Adjusted R-squared:  0.3578
## F-statistic: 41.48 on 11 and 788 DF,  p-value: < 2.2e-16
```

```
pred2 <- predict(linear_model,test)
testerr <- rse(test$quality,pred2 )
```

MSE for linear regression = 0.4091203.

```
#Ridge Regression
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.4
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-1
```

```

x <- model.matrix (quality~.,redwine )[, -1]
y <- redwine$quality
x_train <- x[-test_index,]
y_train <- y[-test_index]
x_test <- x[test_index,]
y_test <- y[test_index]

#sequence of lambdas to be tested
grid = 10^ seq (10, -2, length = 100)

#find the optimal lambda using cross validation
ridge_mod = glmnet (x_train, y_train, alpha = 0, lambda = grid ,
thresh = 1e-12)

set.seed (100)
cv_ridge <- cv.glmnet(x_train, y_train, alpha = 0, lambda = grid)
optimal_lambda <- cv_ridge$lambda.min

#make predictions on the test set
ridge_pred <- predict(ridge_mod, s = optimal_lambda, newx = x_test)
testerr2 <- rse(y_test, ridge_pred)

out = glmnet (x_train, y_train, alpha = 0)
predict (out , type = "coefficients", s = optimal_lambda )

```

```

## 12 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)    20.914828645
## fixed.acidity    0.026536059
## volatile.acidity -1.047451229
## citric.acid      0.031839418
## residual.sugar   0.007182338
## chlorides        -2.173741029
## free.sulfur.dioxide 0.003248019
## total.sulfur.dioxide -0.002995460
## density          -16.643662685
## pH               -0.369405178
## sulphates        0.910337638
## alcohol          0.242468832

```

Optimal value of lambda for ridge regression = 0.070548.
The MSE is 0.4083508.

#Lasso Regression

```

#find the optimal lambda using cross validation
lasso_mod = glmnet (x_train, y_train, alpha = 1, lambda = grid)

set.seed (100)
cv_lasso <- cv.glmnet(x_train, y_train, alpha = 1, lambda = grid)
optimal_lambda <- cv_lasso$lambda.min
optimal_lambda

```

```
## [1] 0.01747528
```

```
#make predictions on the test set
lasso_pred <- predict(lasso_mod, s = optimal_lambda, newx = x_test)
testerr3 <- rse(y_test,lasso_pred)
testerr3
```

```
## [1] 0.4084539
```

```
out=glmnet (x_train,y_train,alpha =1)
predict (out ,type="coefficients",s=optimal_lambda )
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"
##                               1
## (Intercept)                4.251626306
## fixed.acidity               0.009283378
## volatile.acidity           -1.122037373
## citric.acid                 .
## residual.sugar              .
## chlorides                   -1.757534763
## free.sulfur.dioxide         .
## total.sulfur.dioxide       -0.002047706
## density                     .
## pH                          -0.354407573
## sulphates                   0.831489738
## alcohol                    0.267056943
```

Optimal value of lambda for lasso regression = 0.0174753.
The MSE is 0.4084539.

```
#Select best linear model using forward stepwise regression
library (leaps)
```

```
## Warning: package 'leaps' was built under R version 4.0.4
```

```
set.seed(100)
test_index <- sample(c(1:1599),799)
test <- redwine[test_index,]
train <- redwine[-test_index,]
regfit_fwd=regsubsets(quality~.,data=train,method="forward",nvmax=12)
summary(regfit_fwd)
```

```
## Subset selection object
## Call: regsubsets.formula(quality ~ ., data = train, method = "forward",
##       nvmax = 12)
## 11 Variables (and intercept)
##               Forced in Forced out
## fixed.acidity      FALSE      FALSE
## volatile.acidity   FALSE      FALSE
## citric.acid        FALSE      FALSE
```

```

## residual.sugar          FALSE      FALSE
## chlorides               FALSE      FALSE
## free.sulfur.dioxide     FALSE      FALSE
## total.sulfur.dioxide    FALSE      FALSE
## density                 FALSE      FALSE
## pH                      FALSE      FALSE
## sulphates               FALSE      FALSE
## alcohol                 FALSE      FALSE
## 1 subsets of each size up to 11
## Selection Algorithm: forward
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " "*" " " " " "
## 3 ( 1 ) " " "*" " " " " "
## 4 ( 1 ) " " "*" " " " " "
## 5 ( 1 ) " " "*" " " " " "*"
## 6 ( 1 ) " " "*" " " " " "*"
## 7 ( 1 ) " " "*" " " " " "*"
## 8 ( 1 ) " " "*" " " " " "*"
## 9 ( 1 ) " " "*" "*" " " "*"
## 10 ( 1 ) "*" "*" "*" " " "*"
## 11 ( 1 ) "*" "*" "*" "*" "*"
##      free.sulfur.dioxide total.sulfur.dioxide density pH sulphates
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " " "
## 3 ( 1 ) " " " " " " "*"
## 4 ( 1 ) " " "*" " " " "*"
## 5 ( 1 ) " " "*" " " " "*"
## 6 ( 1 ) " " "*" " " "*" "*"
## 7 ( 1 ) "*" "*" " " "*" "*"
## 8 ( 1 ) "*" "*" "*" "*" "*"
## 9 ( 1 ) "*" "*" "*" "*" "*"
## 10 ( 1 ) "*" "*" "*" "*" "*"
## 11 ( 1 ) "*" "*" "*" "*" "*"
##      alcohol
## 1 ( 1 ) "*"
## 2 ( 1 ) "*"
## 3 ( 1 ) "*"
## 4 ( 1 ) "*"
## 5 ( 1 ) "*"
## 6 ( 1 ) "*"
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"
## 11 ( 1 ) "*"

```

```

test_mat=model.matrix(quality~.,data=test)
val_errors =rep(NA,11)

for(i in 1:11){
  # Find the coefficients selected in the models of different sizes
  coefi=coef(regfit_fwd ,id=i)
  #predict the test values

```

```

pred=test_mat[,names(coefi)]%*%coefi
#calculate the errors
val_errors[i]= mean(( test$quality-pred)^2)
}
# number of variables selected
best_model = which.min(val_errors)
best_model

```

```
## [1] 7
```

```

testerr_fwd = val_errors[best_model]
# 7 variables selected over the test set
coef(regfit_fwd,best_model)

```

```

##          (Intercept)      volatile.acidity      chlorides
##      5.031454579      -1.114244843      -2.40007239
## free.sulfur.dioxide total.sulfur.dioxide      pH
##      0.004544875      -0.003472825      -0.605115158
##      sulphates      alcohol
##      0.956816066      0.275876634

```

The MSE is 0.4085524.

```

#Select best linear model using backward stepwise regression
set.seed(100)
test_index <- sample(c(1:1599),799)
test <- redwine[test_index,]
train <- redwine[-test_index,]
regfit_bwd=regsubsets(quality~.,data=train,method="backward",nvmax=12)
summary(regfit_bwd)

```

```

## Subset selection object
## Call: regsubsets.formula(quality ~ ., data = train, method = "backward",
##      nvmax = 12)
## 11 Variables (and intercept)
##
##          Forced in Forced out
## fixed.acidity      FALSE      FALSE
## volatile.acidity      FALSE      FALSE
## citric.acid      FALSE      FALSE
## residual.sugar      FALSE      FALSE
## chlorides      FALSE      FALSE
## free.sulfur.dioxide      FALSE      FALSE
## total.sulfur.dioxide      FALSE      FALSE
## density      FALSE      FALSE
## pH      FALSE      FALSE
## sulphates      FALSE      FALSE
## alcohol      FALSE      FALSE
## 1 subsets of each size up to 11
## Selection Algorithm: backward
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides

```



```

## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " "*" " " " " "
## 3 ( 1 ) " " "*" " " " " "
## 4 ( 1 ) " " "*" " " " " "
## 5 ( 1 ) " " "*" " " " " "*"
## 6 ( 1 ) " " "*" " " " " "*"
## 7 ( 1 ) " " "*" " " " " "*"
## 8 ( 1 ) " " "*" " " " " "*"
## 9 ( 1 ) " " "*" "*" " " "*"
## 10 ( 1 ) "*" "*" "*" " " "*"
## 11 ( 1 ) "*" "*" "*" "*" "*" "*"
##
## free.sulfur.dioxide total.sulfur.dioxide density pH sulphates
## 1 ( 1 ) " " " " " " " "
## 2 ( 1 ) " " " " " " " "
## 3 ( 1 ) " " " " " " "*"
## 4 ( 1 ) " " "*" " " " "*"
## 5 ( 1 ) " " "*" " " " "*"
## 6 ( 1 ) " " "*" " " "*" "*"
## 7 ( 1 ) "*" "*" " " "*" "*"
## 8 ( 1 ) "*" "*" "*" "*" "*"
## 9 ( 1 ) "*" "*" "*" "*" "*"
## 10 ( 1 ) "*" "*" "*" "*" "*"
## 11 ( 1 ) "*" "*" "*" "*" "*"
##
## alcohol
## 1 ( 1 ) "*"
## 2 ( 1 ) "*"
## 3 ( 1 ) "*"
## 4 ( 1 ) "*"
## 5 ( 1 ) "*"
## 6 ( 1 ) "*"
## 7 ( 1 ) "*"
## 8 ( 1 ) "*"
## 9 ( 1 ) "*"
## 10 ( 1 ) "*"
## 11 ( 1 ) "*"

```

```

test_mat=model.matrix(quality~.,data=test)
val_errors =rep(NA,11)

for(i in 1:11){
  # Find the coefficients selected in the models of different sizes
  coefi=coef(regfit_bwd ,id=i)
  #predict the test values
  pred=test_mat[,names(coefi)]%*%coefi
  #calculate the errors
  val_errors[i]= mean((test$quality-pred)^2)
}

best_model = which.min(val_errors)
best_model

```

```
## [1] 7
```

```
testerr_bwd = val_errors[best_model]
# 7 variables selected over the test set

coef(regfit_bwd,best_model)
```

```
##          (Intercept)      volatile.acidity      chlorides
##          5.031454579      -1.114244843      -2.400007239
## free.sulfur.dioxide total.sulfur.dioxide      pH
##          0.004544875      -0.003472825      -0.605115158
##          sulphates      alcohol
##          0.956816066      0.275876634
```

The MSE is 0.4085524.

The test error for all the models are about 0.41.

It appears that the probability that we will correctly predict wine quality is about 0.6.

The same 7 predictors are chosen for lasso, forward and backward regression: volatile acidity, chlorides, free sulfur dioxide, total sulfur dioxide, pH, sulphates and alcohol.

With linear regression, we find 6 variables to be significant : the same as above except free sulfur dioxide.

With ridge regression, all variables are used.

All things considered, we can say that all models (except the ridge model) are fairly similar.

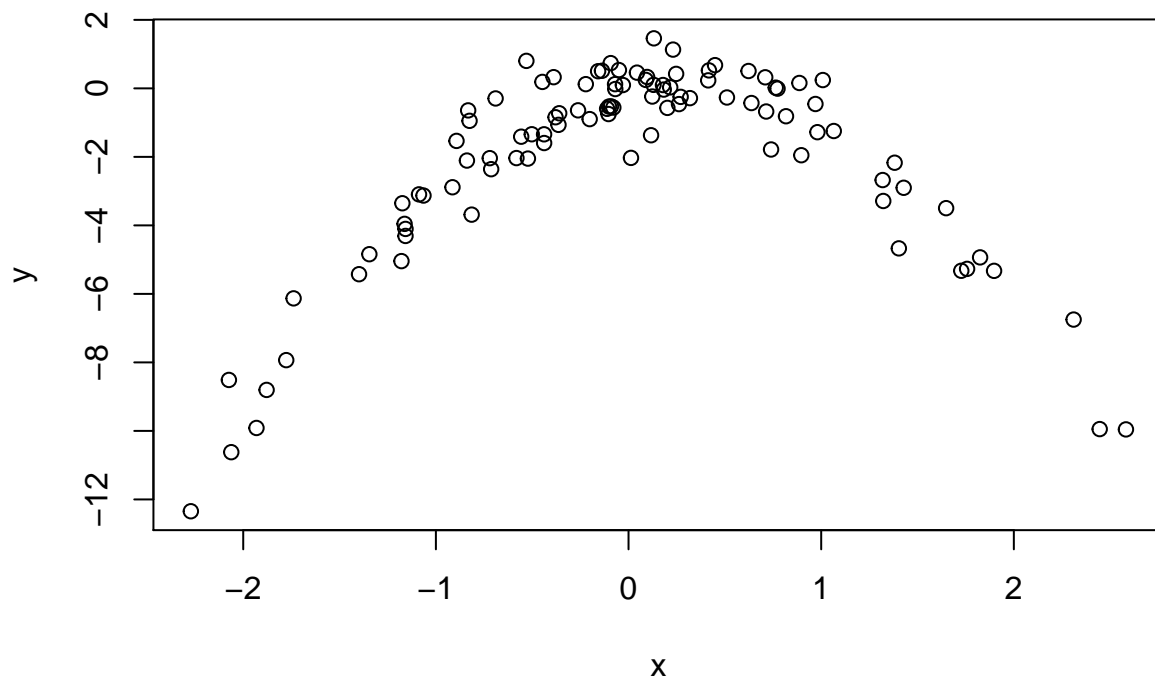
Question 3 - Simulation Study: Cross Validation Method

```
set.seed(100)
x=rnorm(100)
y=x-2*x^2+rnorm(100)
```

n = 100.

$y = (x-2)x^2 + \epsilon$ where $\epsilon \sim N(0, 1)$ p = 2

```
plot(x,y)
```



There appears to be a non linear (quadratic relationship) between x and y.

```
library(boot)
```

```
## Warning: package 'boot' was built under R version 4.0.4
```

```
set.seed(100)
x=rnorm(100)
y=x-2*x^2+rnorm(100)
simulated <- data.frame(x,y)

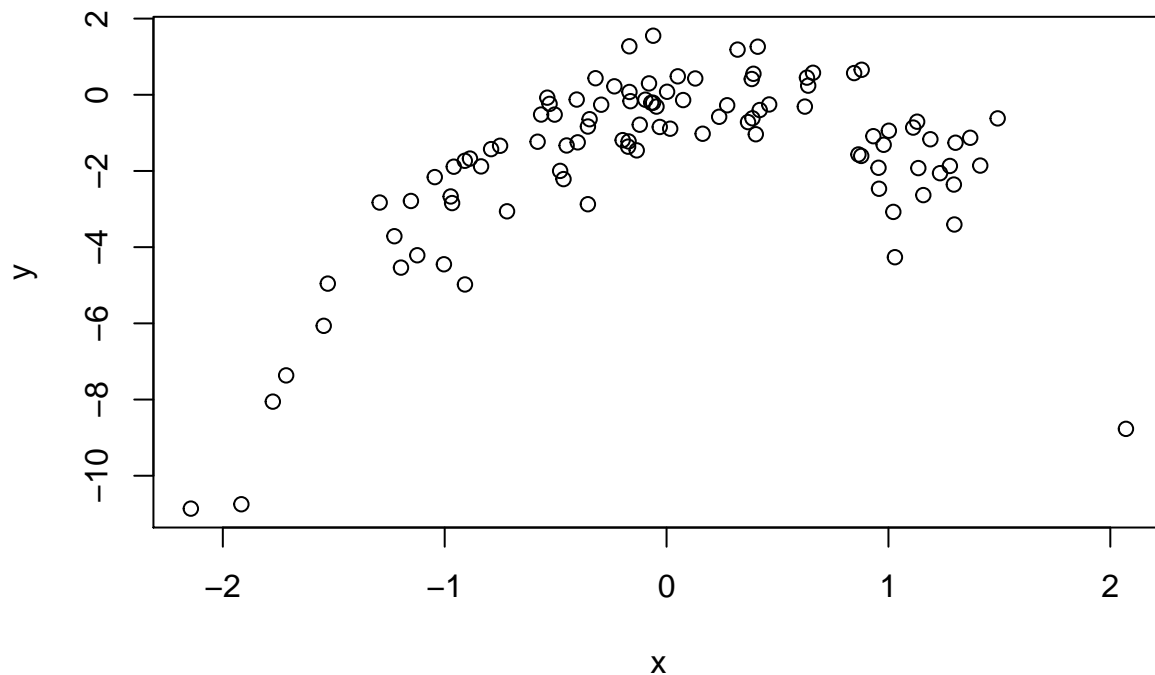
cv_error=rep (0,5)
for (i in 1:5){
  glm_fit=glm(y~poly(x ,i),data=simulated)
  cv_error[i]=cv.glm(simulated,glm_fit)$delta [1]
}
cv_error
```

```
## [1] 9.0606362 0.6511909 0.6665339 0.6671261 0.6744096
```

```
index <- which.min(cv_error)
min_error1 <- cv_error[index]
```

The quadratic model has the smallest LOOCV error : 0.6511909.

```
set.seed(205)
x=rnorm(100)
y=x-2*x^2+rnorm(100)
plot(x,y)
```



```
simulated <- data.frame(x,y)
cv_error=rep(0,5)
for (i in 1:5){
  glm_fit=glm(y~poly(x ,i),data=simulated)
  cv_error[i]=cv.glm(simulated,glm_fit)$delta [1]
}
cv_error
```

```
## [1] 4.755583 1.037305 1.175235 1.427486 3.159778
```

```
index <- which.min(cv_error)
min_error2 <- cv_error[index]
```

Again, the quadratic model has the smallest LOOCV error : 1.0373049.

This is because the scatterplot in both cases is best approximated by a quadratic curve. (even if the true function is cubic).

However even if the quadratic model has the smallest error in both cases, the LOOCV error for the second experiment is larger.

I expected the models to have different errors because each time, the values of x and y will change depending

on the seed we set.

```
set.seed(100)
x=rnorm(100)
y=x-2*x^2+rnorm(100)
simulated <- data.frame(x,y)
```

```
fit1=glm(y~poly(x ,1),data=simulated)
summary(fit1)
```

```
##
## Call:
## glm(formula = y ~ poly(x, 1), data = simulated)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -9.313  -1.212   1.125   1.968   3.439
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.0488     0.2908  -7.045 2.59e-10 ***
## poly(x, 1)    5.5351     2.9080   1.903  0.0599 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 8.456659)
##
##      Null deviance: 859.39  on 99  degrees of freedom
## Residual deviance: 828.75  on 98  degrees of freedom
## AIC: 501.26
##
## Number of Fisher Scoring iterations: 2
```

```
fit2=glm(y~poly(x ,2),data=simulated)
summary(fit2)
```

```
##
## Call:
## glm(formula = y ~ poly(x, 2), data = simulated)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0511  -0.4242  -0.1232   0.5291   1.8763
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.04883     0.07897 -25.945 < 2e-16 ***
## poly(x, 2)1   5.53505     0.78969   7.009 3.2e-10 ***
## poly(x, 2)2 -27.71753     0.78969 -35.099 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for gaussian family taken to be 0.6236173)
##
## Null deviance: 859.389 on 99 degrees of freedom
## Residual deviance: 60.491 on 97 degrees of freedom
## AIC: 241.52
##
## Number of Fisher Scoring iterations: 2
```

```
fit3=glm(y~poly(x ,3),data=simulated)
summary(fit3)
```

```
##
## Call:
## glm(formula = y ~ poly(x, 3), data = simulated)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0027  -0.4533  -0.1187   0.5101   1.8385
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.04883    0.07924  -25.856 < 2e-16 ***
## poly(x, 3)1    5.53505    0.79238   6.985 3.72e-10 ***
## poly(x, 3)2  -27.71753    0.79238 -34.980 < 2e-16 ***
## poly(x, 3)3    0.46381    0.79238   0.585    0.56
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.6278725)
##
## Null deviance: 859.389 on 99 degrees of freedom
## Residual deviance: 60.276 on 96 degrees of freedom
## AIC: 243.16
##
## Number of Fisher Scoring iterations: 2
```

```
fit4=glm(y~poly(x ,4),data=simulated)
summary(fit4)
```

```
##
## Call:
## glm(formula = y ~ poly(x, 4), data = simulated)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0603  -0.4947  -0.1133   0.5593   1.8436
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -2.04883    0.07883  -25.991 < 2e-16 ***
## poly(x, 4)1    5.53505    0.78829   7.022 3.26e-10 ***
## poly(x, 4)2  -27.71753    0.78829 -35.162 < 2e-16 ***
## poly(x, 4)3    0.46381    0.78829   0.588    0.558
```

```
## poly(x, 4) 1.11467 0.78829 1.414 0.161
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.6214028)
##
## Null deviance: 859.389 on 99 degrees of freedom
## Residual deviance: 59.033 on 95 degrees of freedom
## AIC: 243.08
##
## Number of Fisher Scoring iterations: 2
```

As illustrated by the summaries, the AIC value for the quadratic model is the smallest (as suggested by the cross-validation method).

For the linear model, the coefficient for β_1 is not significant.

For the quadratic model, both β_1 and β_2 are significant.

For the cubic model, both β_1 and β_2 are significant but β_3 is not.

For the quartic model, β_1 and β_2 are significant but β_3 and β_4 are not.

These observations agree with the results of the LOOCV, which selected the quadratic model.

Question 4 - Simulation Study: Screening, Stepwise Selection and ROC Curves

```
gwas <- read.csv("GWAS.CSV")
gwas <- gwas[,-1]
#Convert all the columns to factors except V1
set.seed(100)
train_index <- sample( c(1:3000), 1500)
train <- gwas[train_index,]
test <- gwas[-train_index,]
logistic <- glm(V1~.,data=train,family="binomial")
#Significant attributes
significant <- summary(logistic)$coef[,4][summary(logistic)$coef[,4] < 0.05]
prob <- predict(logistic,test,type="response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type == :
## prediction from a rank-deficient fit may be misleading
```

```
pred <- ifelse(prob >= 0.5, 1, 0)
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'lattice'
```

```
## The following object is masked from 'package:boot':
```

```
##
```

```
## melanoma
```

```
## Loading required package: ggplot2
```

```
confusionMatrix(as.factor(pred),as.factor(test$V1))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 646 336
##           1 324 194
##
##           Accuracy : 0.56
##           95% CI : (0.5344, 0.5853)
##       No Information Rate : 0.6467
##       P-Value [Acc > NIR] : 1.0000
##
##           Kappa : 0.0322
##
##  Mcnemar's Test P-Value : 0.6685
##
##           Sensitivity : 0.6660
##           Specificity : 0.3660
##       Pos Pred Value : 0.6578
##       Neg Pred Value : 0.3745
##           Prevalence : 0.6467
##       Detection Rate : 0.4307
##       Detection Prevalence : 0.6547
##       Balanced Accuracy : 0.5160
##
##       'Positive' Class : 0
##
```

/ On test data set, the model predicts correctly 56% of the time./ It is not a good model./

```
#Ridge Regression
library(glmnet)
x <- model.matrix (V1~.,train )[, -1]
y <- train$V1
x_test<- model.matrix (V1~.,test )[, -1]
y_test <- test$V1

#sequence of lambdas to be tested
grid =10^ seq (10,-2, length =100)

#find the optimal lambda using cross validation

set.seed (100)
cv_ridge <- cv.glmnet(x, y, alpha = 0, lambda = grid,family="binomial")
optimal_lambda <- cv_ridge$lambda.min

#make predictions on the test set
ridge_prob<- predict(cv_ridge, s = optimal_lambda, newx = x_test,type="response")
```



```
ridge_pred <- ifelse(ridge_prob >= 0.5, 1, 0)
confusionMatrix(as.factor(y), as.factor(ridge_pred))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 964    0
##           1 534    2
##
##           Accuracy : 0.644
##           95% CI : (0.6192, 0.6683)
##           No Information Rate : 0.9987
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0048
##
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.643525
##           Specificity : 1.000000
##           Pos Pred Value : 1.000000
##           Neg Pred Value : 0.003731
##           Prevalence : 0.998667
##           Detection Rate : 0.642667
##           Detection Prevalence : 0.642667
##           Balanced Accuracy : 0.821762
##
##           'Positive' Class : 0
##
```

The ridge logistic regression model is 64.4% accurate.

#Lasso Regression

#find the optimal lambda using cross validation

```
set.seed(100)
cv_lasso <- cv.glmnet(x, y, alpha = 1, lambda = grid, family="binomial")
optimal_lambda <- cv_lasso$lambda.min
```

#make predictions on the test set

```
lasso_prob <- predict(cv_lasso, s = optimal_lambda, newx = x_test, type="response")
lasso_pred <- ifelse(lasso_prob >= 0.5, 1, 0)
table(lasso_pred, y)
```

```
##           y
## lasso_pred  0    1
##           0 964 536
```

```
mean(lasso_pred==y)
```

```
## [1] 0.6426667
```

The lasso logistic regression model is 64.27 % accurate but it classifies all the test data as 0.

Summary:

Model Accuracy

Logistic 56%

Ridge Logistic 64.4%

Lasso Logistic 64.3%

It appears that the penalised logistic regression models (Lasso and Ridge) work better than the traditional logistic regression model.

```
library(pROC)
```

```
## Warning: package 'pROC' was built under R version 4.0.4
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
par(pty = "s")
```

```
#plot ROC curve for ridge logistic regression
```

```
roc_info_ridge <- roc(response=test$V1, predictor=ridge_prob , plot= TRUE, legacy.axes=TRUE,col="red")
```

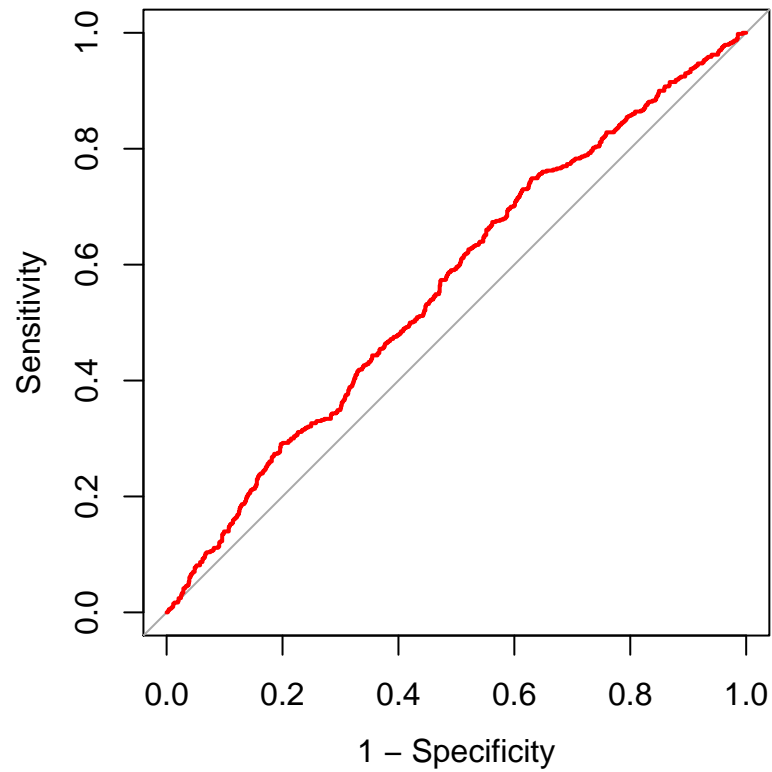
```
## Setting levels: control = 0, case = 1
```

```
## Warning in roc.default(response = test$V1, predictor = ridge_prob, plot
```

```
## = TRUE, : Deprecated use a matrix as predictor. Unexpected results may be
```

```
## produced, please pass a numeric vector.
```

```
## Setting direction: controls < cases
```

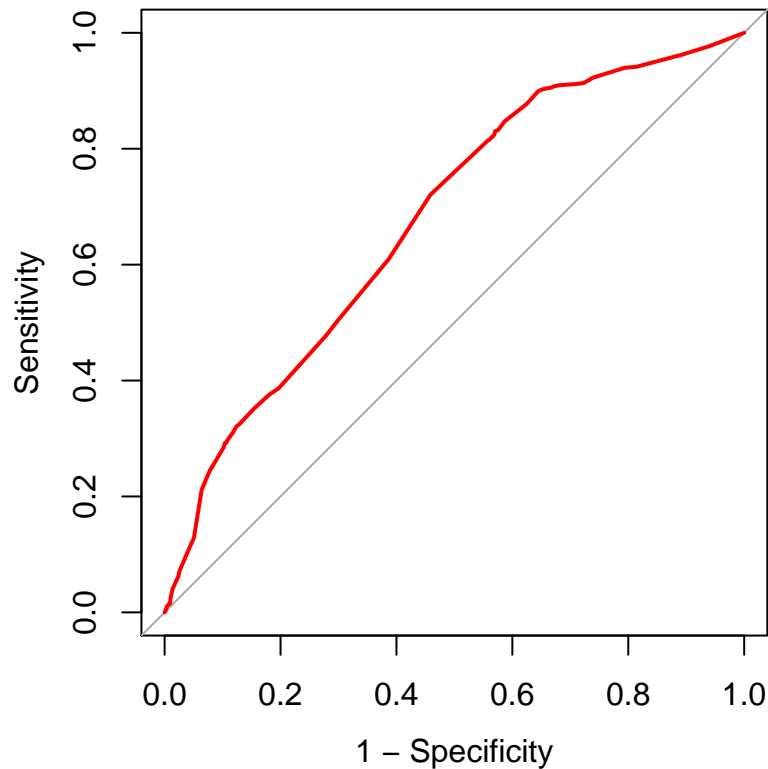


```
#plot ROC curve for lasso logistic regression
roc_info_lasso <- roc(response=test$V1, predictor=lasso_prob , plot= TRUE, legacy.axes=TRUE,col="red")

## Setting levels: control = 0, case = 1

## Warning in roc.default(response = test$V1, predictor = lasso_prob, plot
## = TRUE, : Deprecated use a matrix as predictor. Unexpected results may be
## produced, please pass a numeric vector.

## Setting direction: controls < cases
```



```
roc.df <- data.frame(tp=roc_info_lasso$sensitivities*100,
fpp=(1-roc_info_lasso$specificities)*100,thresholds=roc_info_lasso$thresholds)
```

```
threshold <- seq(0,1,by=0.001)
accuracy <- rep(0,length(threshold))
for(i in 1:length(threshold)){
  lasso_new_pred <- ifelse(lasso_prob >= threshold[i], 1, 0)
  accuracy[i] <- mean(lasso_new_pred==y)
}
index <- which.max(accuracy)
accuracy[index]
```

```
## [1] 0.6426667
```

```
threshold[index]
```

```
## [1] 0.498
```

It seems that for a threshold above or equal to 0.498, the highest accuracy = 0.6426667 remains constant. This is because the lasso model at a threshold of 0.498 classifies all observations as 0. Thus, increasing the threshold above 0.50 does not change the accuracy but decreasing the threshold leads to a decrease in accuracy.

```
library(coefplot)
```

```
## Warning: package 'coefplot' was built under R version 4.0.4
```

```
extract.coef(cv.glmnet(x, y, alpha = 1, lambda = grid,family="binomial"))
```

```
##              Value SE Coefficient
## (Intercept) -0.45289612 NA (Intercept)
## V65          0.29290046 NA          V65
## V80          0.10746436 NA          V80
## V300         -0.66302796 NA          V300
## V316          0.04070697 NA          V316
## V337          0.02334652 NA          V337
## V408         -0.03933016 NA          V408
```

Decision procedure:

Collect data about alleles V65, V80, V300, V316, V337 and V408 for the population.

We choose these alleles because these are the only variables that have non zero coefficients in the lasso model.

We predict the chances of having the disease using the lasso model.

We want the number of false negatives to be as small as possible (as we don't want a patient having the disease to be classified as not having the disease)

One course of action would be to establish an accepted level of false negatives and try to reduce the number of false positives (i.e aim for a high sensibility), by choosing an appropriate threshold based on the ROC graph.

We then send these people for screening

Question 7 - Empirical Study: LASSO vs. Regression Workflow

Bhavika Sewpal - 300089940

4/10/2021

```
train <- read.csv("ListingsTrain.csv")

#Cleaning up for train dataset

#Removing columns neighbourhood_group_cleansed and bathroom (contain only NAs)
train <- train[,c(-13,-19)]

#Converting categorical variables (true/false) to 0 and 1.
train$host_is_superhost <- factor(train$host_is_superhost, levels = c("t","f"), labels=c(1,0))

train$host_has_profile_pic <- factor(train$host_has_profile_pic, levels = c("t","f"), labels=c(1,0))
train$host_identity_verified <- factor(train$host_identity_verified, levels = c("t","f"), labels=c(1,0))
train$has_availability<- factor(train$has_availability, levels = c("t","f"), labels=c(1,0))
train$instant_bookable <- factor(train$instant_bookable, levels = c("t","f"), labels=c(1,0))

#Remove 72 rows
index <- is.na(train$review_scores_checkin)
train <- train[(!index),]

#Remove dollar sign from price column
train$price = as.numeric(gsub("\\$", "", train$price))

## Warning: NAs introduced by coercion

#find out the NAs in price column and replace them with the average value of price
index1 <- which(is.na(train$price))
median_price <- 89
train[index1,21] = median_price

#Remove percentage sign in host_response_rate column
train$host_response_rate = as.numeric(gsub("%", "", train$host_response_rate))

## Warning: NAs introduced by coercion

#find out the NAs in the response rate column and replace them with the average value of response rate
index2 <- which(is.na(train$host_response_rate))
avg_response <- 92.5
train[index2,4] = avg_response
```

```
#Remove percentage sign in host_acceptance_rate column
train$host_acceptance_rate = as.numeric(gsub("\\%", "", train$host_acceptance_rate))
```

```
## Warning: NAs introduced by coercion
```

```
#find out the NAs in the acceptance rate column and replace them with the average value of acceptance r
index3 <- which(is.na(train$host_acceptance_rate))
avg_acceptance <- 76.87
train[index3,5] = avg_acceptance
```

```
#Replace missing values in host_response_time with more frequent category:within an hour
index4 <- which((train$host_response_time=="N/A"))
train[index4,3] <- "within an hour"
```

```
index5 <- which(is.na(train$bedrooms))
avg_bedrooms <- 1.7
train[index5,18] <- 1.7
```

```
#Converting to categorical variables to factors
train$room_type <- factor(train$room_type)
train$host_response_time <- factor(train$host_response_time)
train$host_neighbourhood<- factor(train$host_neighbourhood)
train$neighbourhood_cleansed<- factor(train$neighbourhood_cleansed)
train$property_type <- factor(train$property_type)
```

```
summary(train)
```

```
## description      neighborhood_overview      host_response_time
## Length:928      Length:928      a few days or more: 23
## Class :character Class :character      within a day      : 79
## Mode  :character Mode  :character      within a few hours:132
##                                     within an hour    :694
##
##
##
## host_response_rate host_acceptance_rate host_is_superhost
## Min.   : 0.00      Min.   : 0.00      1:361
## 1st Qu.: 92.50      1st Qu.: 76.40      0:567
## Median : 99.00      Median : 80.00
## Mean   : 92.78      Mean    : 77.62
## 3rd Qu.:100.00      3rd Qu.: 97.25
## Max.   :100.00      Max.    :100.00
##
##                host_neighbourhood host_listings_count
##                :827      Min.   : 0.000
## Sandy Hill      : 24      1st Qu.: 1.000
## Byward Market - Parliament Hill: 11      Median : 2.000
## Centretown      : 11      Mean    : 5.413
## Downtown        : 10      3rd Qu.: 3.000
## Lower Town      : 6       Max.    :272.000
## (Other)         : 39
```

```

## host_total_listings_count host_has_profile_pic host_identity_verified
## Min. : 0.000 1:927 1:781
## 1st Qu.: 1.000 0: 1 0:147
## Median : 2.000
## Mean : 5.413
## 3rd Qu.: 3.000
## Max. :272.000
##
## neighbourhood_cleansed latitude longitude
## Rideau-Vanier :176 Min. :45.13 Min. : -76.22
## Somerset :132 1st Qu.:45.37 1st Qu.: -75.73
## Kitchissippi : 95 Median :45.41 Median : -75.69
## Capital : 78 Mean :45.39 Mean : -75.70
## River : 47 3rd Qu.:45.43 3rd Qu.: -75.67
## Rideau-Rockcliffe: 41 Max. :45.51 Max. : -75.35
## (Other) :359
## property_type room_type accommodates
## Entire apartment :240 Entire home/apt:583 Min. : 1.000
## Private room in house:201 Hotel room : 1 1st Qu.: 2.000
## Entire house :152 Private room :338 Median : 3.000
## Entire guest suite : 49 Shared room : 6 Mean : 3.642
## Entire condominium : 46 3rd Qu.: 5.000
## Entire townhouse : 45 Max. :16.000
## (Other) :195
## bedrooms beds amenities price
## Min. :1.0 Min. : 0.000 Length:928 Min. : 0.0
## 1st Qu.:1.0 1st Qu.: 1.000 Class :character 1st Qu.: 60.0
## Median :1.0 Median : 1.000 Mode :character Median : 89.0
## Mean :1.7 Mean : 1.928 Mean : 133.9
## 3rd Qu.:2.0 3rd Qu.: 3.000 3rd Qu.: 125.0
## Max. :6.0 Max. :12.000 Max. :9800.0
##
## minimum_nights maximum_nights has_availability number_of_reviews
## Min. : 1.00 Min. :1.000e+00 1:913 Min. : 1.00
## 1st Qu.: 1.00 1st Qu.:5.000e+01 0: 15 1st Qu.: 12.00
## Median : 2.00 Median :1.125e+03 Median : 39.00
## Mean : 11.93 Mean :2.315e+06 Mean : 73.13
## 3rd Qu.: 5.00 3rd Qu.:1.125e+03 3rd Qu.: 98.25
## Max. :1000.00 Max. :2.147e+09 Max. :553.00
##
## number_of_reviews_ltm number_of_reviews_l30d first_review
## Min. : 0.000 Min. :0.0000 Length:928
## 1st Qu.: 0.000 1st Qu.:0.0000 Class :character
## Median : 1.000 Median :0.0000 Mode :character
## Mean : 6.534 Mean :0.2263
## 3rd Qu.: 7.000 3rd Qu.:0.0000
## Max. :112.000 Max. :6.0000
##
## last_review review_scores_rating review_scores_accuracy
## Length:928 Min. : 20.00 Min. : 2.000
## Class :character 1st Qu.: 94.00 1st Qu.:10.000
## Mode :character Median : 97.00 Median :10.000
## Mean : 95.38 Mean : 9.749
## 3rd Qu.: 99.00 3rd Qu.:10.000

```



```
##           Max.      :100.00           Max.      :10.000
##
## review_scores_cleanliness review_scores_checkin review_scores_communication
## Min.      : 2.00           Min.      : 2.000           Min.      : 2.000
## 1st Qu.: 9.00           1st Qu.:10.000           1st Qu.:10.000
## Median :10.00           Median :10.000           Median :10.000
## Mean      : 9.58           Mean      : 9.894           Mean      : 9.874
## 3rd Qu.:10.00           3rd Qu.:10.000           3rd Qu.:10.000
## Max.      :10.00           Max.      :10.000           Max.      :10.000
##
## review_scores_location review_scores_value instant_bookable
## Min.      : 2.000           Min.      : 2.000           1:191
## 1st Qu.:10.000           1st Qu.: 9.000           0:737
## Median :10.000           Median :10.000
## Mean      : 9.706           Mean      : 9.564
## 3rd Qu.:10.000           3rd Qu.:10.000
## Max.      :10.000           Max.      :10.000
##
## calculated_host_listings_count calculated_host_listings_count_entire_homes
## Min.      : 1.000           Min.      : 0.000
## 1st Qu.: 1.000           1st Qu.: 0.000
## Median : 1.000           Median : 1.000
## Mean      : 3.454           Mean      : 2.505
## 3rd Qu.: 3.000           3rd Qu.: 2.000
## Max.      :69.000           Max.      :69.000
##
## calculated_host_listings_count_private_rooms reviews_per_month
## Min.      : 0.0000           Min.      :0.010
## 1st Qu.: 0.0000           1st Qu.:0.270
## Median : 0.0000           Median :0.800
## Mean      : 0.9332           Mean      :1.509
## 3rd Qu.: 1.0000           3rd Qu.:2.060
## Max.      :12.0000           Max.      :9.810
##
```

```
#Cleaning up optimization dataset
```

```
#Removing columns neighbourhood_group_cleaned and bathroom (contain only NAs)
```

```
opt <- read.csv("ListingsOptimization.csv")
```

```
opt <- opt[,c(-13,-19)]
```

```
#Converting categorical variables (true/false) to 0 and 1.
```

```
opt$host_is_superhost <- factor(opt$host_is_superhost, levels = c("t","f"), labels=c(1,0))
```

```
opt$host_has_profile_pic <- factor(opt$host_has_profile_pic, levels = c("t","f"), labels=c(1,0))
```

```
opt$host_identity_verified <- factor(opt$host_identity_verified, levels = c("t","f"), labels=c(1,0))
```

```
opt$has_availability <- factor(opt$has_availability, levels = c("t","f"), labels=c(1,0))
```

```
opt$instant_bookable <- factor(opt$instant_bookable, levels = c("t","f"), labels=c(1,0))
```

```
#Remove 72 rows
```

```
index <- is.na(opt$review_scores_checkin)
```

```
opt <- opt[(!index),]
```

```
#Remove dollar sign from price column
```

```
opt$price = as.numeric(gsub("\\$", "", opt$price))
```

```
## Warning: NAs introduced by coercion
```

```
#find out the NAs in price column and replace them with the average value of price  
index1 <- which(is.na(opt$price))  
median_price <- 80  
opt[index1,21] <- median_price
```

```
#Remove percentage sign in host_response_rate column  
opt$host_response_rate = as.numeric(gsub("%", "", opt$host_response_rate))
```

```
## Warning: NAs introduced by coercion
```

```
#find out the NAs in the response rate column and replace them with the average value of response rate  
index2 <- which(is.na(opt$host_response_rate))  
avg_response <- 92.76  
opt[index2,4] <- avg_response
```

```
#Remove percentage sign in host_acceptance_rate column  
opt$host_acceptance_rate = as.numeric(gsub("%", "", opt$host_acceptance_rate))
```

```
## Warning: NAs introduced by coercion
```

```
#find out the NAs in the acceptance rate column and replace them with the average value of acceptance rate  
index3 <- which(is.na(opt$host_acceptance_rate))  
avg_acceptance <- 82.98  
opt[index3,5] <- avg_acceptance
```

```
#Replace missing values in host_response_time with more frequent category:within an hour  
index4 <- which((opt$host_response_time=="N/A"))  
opt[index4,3] <- "within an hour"
```

```
#Converting to categorical variables to factors  
opt$room_type <- factor(opt$room_type)  
opt$host_response_time <- factor(opt$host_response_time)  
opt$host_neighbourhood <- factor(opt$host_neighbourhood)  
opt$neighbourhood_cleansed <- factor(opt$neighbourhood_cleansed)  
opt$property_type <- factor(opt$property_type)
```

```
#dealing with NAs  
index5 <- which(is.na(opt$host_is_superhost))  
opt[index5,6] = 0
```

```
index6 <- which(is.na(opt$host_listings_count))  
opt[index6,8] = 2
```

```
index7 <- which(is.na(opt$host_total_listings_count))  
opt[index7,9] = 2
```

```
index8 <- which(is.na(opt$host_has_profile_pic))  
opt[index8,10] = 1
```

```
index9<- which(is.na(opt$host_identity_verified))
opt[index9,11] = 1
```

```
index10 <- which(is.na(opt$bedrooms))
opt[index10,18] = 1
```

```
summary(opt)
```

```
## description      neighborhood_overview      host_response_time
## Length:848      Length:848                  : 1
## Class :character Class :character      a few days or more: 16
## Mode  :character Mode  :character      within a day      : 50
##                                     within a few hours: 84
##                                     within an hour    :697
##
##
##
## host_response_rate host_acceptance_rate host_is_superhost
## Min.   : 0.00      Min.   : 0.00      1:355
## 1st Qu.: 92.76      1st Qu.: 82.98      0:493
## Median :100.00      Median : 91.00
## Mean   : 94.12      Mean   : 84.91
## 3rd Qu.:100.00      3rd Qu.: 99.00
## Max.   :100.00      Max.   :100.00
##
##
##                host_neighbourhood host_listings_count
##                :710      Min.   : 0.00
## Byward Market - Parliament Hill: 22      1st Qu.: 1.00
## Sandy Hill                      : 18      Median : 2.00
## Centretown                      : 13      Mean   : 7.14
## Downtown                       : 13      3rd Qu.: 4.00
## Centretown West                 : 8       Max.   :272.00
## (Other)                         : 64
## host_total_listings_count host_has_profile_pic host_identity_verified
## Min.   : 0.00      1:841      1:667
## 1st Qu.: 1.00      0: 7      0:181
## Median : 2.00
## Mean   : 7.14
## 3rd Qu.: 4.00
## Max.   :272.00
##
## neighbourhood_cleansed latitude longitude
## Rideau-Vanier:165      Min.   :44.99 Min.   : -76.11
## Somerset :111      1st Qu.:45.35 1st Qu.: -75.74
## Kanata North : 53      Median :45.40 Median : -75.69
## River : 49      Mean   :45.38 Mean   : -75.71
## Capital : 48      3rd Qu.:45.43 3rd Qu.: -75.67
## Kitchissippi : 42      Max.   :45.52 Max.   : -75.39
## (Other) :380
##
## property_type room_type accommodates
## Entire apartment :196 Entire home/apt:500 Min.   : 1.000
## Private room in house :163 Hotel room : 5 1st Qu.: 2.000
## Entire house :123 Private room :337 Median : 3.000
```

```

## Private room in townhouse: 78 Shared room : 6 Mean : 3.586
## Entire guest suite : 55 3rd Qu.: 5.000
## Entire condominium : 41 Max. :16.000
## (Other) :192
## bedrooms beds amenities price
## Min. :1.000 Min. : 0.000 Length:848 Min. : 20.00
## 1st Qu.:1.000 1st Qu.: 1.000 Class :character 1st Qu.: 53.00
## Median :1.000 Median : 1.000 Mode :character Median : 80.00
## Mean :1.605 Mean : 1.909 Mean : 96.29
## 3rd Qu.:2.000 3rd Qu.: 2.000 3rd Qu.:115.00
## Max. :8.000 Max. :16.000 Max. :868.00
##
## minimum_nights maximum_nights has_availability number_of_reviews
## Min. : 1.000 Min. : 1.0 1:842 Min. : 1.00
## 1st Qu.: 1.000 1st Qu.: 31.0 0: 6 1st Qu.: 5.00
## Median : 2.000 Median :1125.0 Median : 17.00
## Mean : 8.888 Mean : 643.1 Mean : 32.23
## 3rd Qu.: 4.000 3rd Qu.:1125.0 3rd Qu.: 44.25
## Max. :1000.000 Max. :1125.0 Max. :279.00
##
## number_of_reviews_ltm number_of_reviews_l30d first_review
## Min. : 0.0 Min. :0.0000 Length:848
## 1st Qu.: 0.0 1st Qu.:0.0000 Class :character
## Median : 3.0 Median :0.0000 Mode :character
## Mean : 9.3 Mean :0.3396
## 3rd Qu.: 11.0 3rd Qu.:0.0000
## Max. :118.0 Max. :9.0000
##
## last_review review_scores_rating review_scores_accuracy
## Length:848 Min. : 20.00 Min. : 2.000
## Class :character 1st Qu.: 93.00 1st Qu.:10.000
## Mode :character Median : 97.00 Median :10.000
## Mean : 94.35 Mean : 9.611
## 3rd Qu.:100.00 3rd Qu.:10.000
## Max. :100.00 Max. :10.000
##
## review_scores_cleanliness review_scores_checkin review_scores_communication
## Min. : 2.000 Min. : 2.000 Min. : 2.000
## 1st Qu.: 9.000 1st Qu.:10.000 1st Qu.:10.000
## Median :10.000 Median :10.000 Median :10.000
## Mean : 9.495 Mean : 9.789 Mean : 9.739
## 3rd Qu.:10.000 3rd Qu.:10.000 3rd Qu.:10.000
## Max. :10.000 Max. :10.000 Max. :10.000
##
## review_scores_location review_scores_value instant_bookable
## Min. : 2.000 Min. : 2.000 1:303
## 1st Qu.:10.000 1st Qu.: 9.000 0:545
## Median :10.000 Median :10.000
## Mean : 9.678 Mean : 9.441
## 3rd Qu.:10.000 3rd Qu.:10.000
## Max. :10.000 Max. :10.000
##
## calculated_host_listings_count calculated_host_listings_count_entire_homes
## Min. : 1.000 Min. : 0.000

```

```
## 1st Qu.: 1.000          1st Qu.: 0.000
## Median : 2.000          Median : 1.000
## Mean   : 4.421          Mean    : 3.179
## 3rd Qu.: 4.000          3rd Qu.: 2.000
## Max.   :69.000          Max.    :69.000
##
## calculated_host_listings_count_private_rooms reviews_per_month
## Min.    : 0.000          Min.    :0.0300
## 1st Qu.: 0.000          1st Qu.:0.2575
## Median : 0.000          Median :0.8000
## Mean    : 1.202          Mean    :1.4326
## 3rd Qu.: 2.000          3rd Qu.:2.0400
## Max.    :12.000          Max.    :9.8800
##
```

#Cleaning up for test dataset

#Removing columns neighbourhood_group_cleansed and bathroom (contain only NAs)

```
test <- read.csv("ListingsTest.csv")
test <- test[,c(-13,-19)]
```

#Converting categorical variables (true/false) to 0 and 1.

```
test$host_is_superhost <- factor(test$host_is_superhost, levels = c("t","f"), labels=c(1,0))

test$host_has_profile_pic <- factor(test$host_has_profile_pic, levels = c("t","f"), labels=c(1,0))
test$host_identity_verified <- factor(test$host_identity_verified, levels = c("t","f"), labels=c(1,0))
test$has_availability <- factor(test$has_availability, levels = c("t","f"), labels=c(1,0))
test$instant_bookable <- factor(test$instant_bookable, levels = c("t","f"), labels=c(1,0))
```

#Remove 400 rows

```
index <- is.na(test$review_scores_checkin)
test <- test[(!index),]
```

#Remove dollar sign from price column

```
test$price = as.numeric(gsub("\\$", "", test$price))
```

#find out the NAs in price column and replace them with the average value of price

```
index1 <- which(is.na(test$price))
median_price <- 80
test[index1,21] <- median_price
```

#Remove percentage sign in host_response_rate column

```
test$host_response_rate = as.numeric(gsub("%", "", test$host_response_rate))
```

```
## Warning: NAs introduced by coercion
```

#find out the NAs in the response rate column and replace them with the average value of response rate

```
index2 <- which(is.na(test$host_response_rate))
avg_response <- 94.34
test[index2,4] <- avg_response
```

#Remove percentage sign in host_acceptance_rate column

```
test$host_acceptance_rate = as.numeric(gsub("%", "", test$host_acceptance_rate))
```

```
## Warning: NAs introduced by coercion
```

```
#find out the NAs in the acceptance rate column and replace them with the average value of acceptance r
index3 <- which(is.na(test$host_acceptance_rate))
avg_acceptance <- 88.18
test[index3,5] <- avg_acceptance

#Replace missing values in host_response_time with more frequent category:within an hour
index4 <- which((test$host_response_time=="N/A"))
test[index4,3] <- "within an hour"

#Converting to categorical variables to factors
test$room_type <- factor(test$room_type)
test$host_response_time <- factor(test$host_response_time)
test$host_neighbourhood<- factor(test$host_neighbourhood)
test$neighbourhood_cleansed<- factor(test$neighbourhood_cleansed)
test$property_type <- factor(test$property_type)

index5 <- which(is.na(test$bedrooms))
test[index5,18] = 1

index6 <- which(is.na(test$beds))
test[index6,19] = 1
summary(test)
```

```
## description      neighborhood_overview      host_response_time
## Length:404      Length:404      a few days or more: 5
## Class :character Class :character      within a day      : 22
## Mode :character Mode :character      within a few hours: 48
##                                     within an hour    :329
##
##
##
## host_response_rate host_acceptance_rate host_is_superhost
## Min. : 0.00      Min. : 0.00      1:113
## 1st Qu.: 99.00    1st Qu.: 89.00    0:291
## Median :100.00    Median : 95.00
## Mean : 96.41      Mean : 90.57
## 3rd Qu.:100.00    3rd Qu.:100.00
## Max. :100.00      Max. :100.00
##
##                host_neighbourhood host_listings_count
##                :153      Min. : 0.000
## Downtown      : 41      1st Qu.: 0.000
## Byward Market - Parliament Hill: 28      Median : 1.000
## Centretown West: 27      Mean : 9.077
## Centretown     : 20      3rd Qu.: 4.000
## Sandy Hill     : 17      Max. :272.000
## (Other)        :118
##
## host_total_listings_count host_has_profile_pic host_identity_verified
## Min. : 0.000      1:404      1:334
## 1st Qu.: 0.000      0: 0      0: 70
## Median : 1.000
```

```

## Mean      : 9.077
## 3rd Qu.: 4.000
## Max.      :272.000
##
##      neighbourhood_cleansed    latitude    longitude
## Rideau-Vanier      :111      Min.    :45.05    Min.    : -76.00
## Somerset          : 58      1st Qu.:45.36    1st Qu.: -75.73
## Kitchissippi       : 26      Median  :45.41    Median  : -75.69
## Rideau-Rockcliffe: 24      Mean     :45.39    Mean     : -75.70
## Capital            : 22      3rd Qu.:45.43    3rd Qu.: -75.67
## College            : 16      Max.     :45.49    Max.     : -75.46
## (Other)            :147
##
##      property_type    room_type    accommodates
## Entire apartment      :159    Entire home/apt:270    Min.    : 1.000
## Entire house          : 44    Private room   :132    1st Qu.: 2.000
## Private room in house : 41    Shared room    : 2    Median  : 2.000
## Private room in townhouse: 38    Mean          : 3.391
## Entire condominium    : 25    3rd Qu.: 4.000
## Entire guest suite     : 20    Max.         :16.000
## (Other)                : 77
##
##      bedrooms    beds    amenities    price
## Min.    :1.000    Min.    :0.00    Length:404    Min.    : 18.00
## 1st Qu.:1.000    1st Qu.:1.00    Class :character    1st Qu.: 52.75
## Median :1.000    Median :1.00    Mode  :character    Median  : 78.50
## Mean    :1.562    Mean    :1.79    Mean    : 89.74
## 3rd Qu.:2.000    3rd Qu.:2.00    3rd Qu.:106.75
## Max.    :9.000    Max.    :9.00    Max.    :700.00
##
##      minimum_nights    maximum_nights    has_availability    number_of_reviews
## Min.    : 1.00    Min.    : 1.0    1:403    Min.    : 1.000
## 1st Qu.: 1.00    1st Qu.: 60.0    0: 1    1st Qu.: 2.000
## Median  : 2.00    Median  : 682.0    Median  : 4.000
## Mean    : 6.53    Mean    : 643.9    Mean    : 8.342
## 3rd Qu.: 4.00    3rd Qu.:1125.0    3rd Qu.:11.000
## Max.    :1000.00    Max.    :1125.0    Max.    :81.000
##
##      number_of_reviews_ltm    number_of_reviews_l30d    first_review
## Min.    : 1.000    Min.    :0.0000    Length:404
## 1st Qu.: 2.000    1st Qu.:0.0000    Class :character
## Median  : 4.000    Median  :0.0000    Mode  :character
## Mean    : 8.104    Mean    :0.8639
## 3rd Qu.:10.250    3rd Qu.:1.0000
## Max.    :69.000    Max.    :8.0000
##
##      last_review    review_scores_rating    review_scores_accuracy
## Length:404    Min.    : 20.00    Min.    : 2.0
## Class :character    1st Qu.: 92.00    1st Qu.: 9.0
## Mode  :character    Median  : 98.00    Median  :10.0
## Mean    : 93.07    Mean    : 9.5
## 3rd Qu.:100.00    3rd Qu.:10.0
## Max.    :100.00    Max.    :10.0
##
##      review_scores_cleanliness    review_scores_checkin    review_scores_communication
## Min.    : 2.000    Min.    : 2.000    Min.    : 2.000

```

```
## 1st Qu.: 9.000          1st Qu.:10.000          1st Qu.:10.000
## Median :10.000          Median :10.000          Median :10.000
## Mean   : 9.413          Mean    : 9.723          Mean    : 9.678
## 3rd Qu.:10.000          3rd Qu.:10.000          3rd Qu.:10.000
## Max.    :10.000          Max.     :10.000          Max.     :10.000
##
## review_scores_location review_scores_value instant_bookable
## Min.    : 2.000          Min.     : 2.000          1:162
## 1st Qu.:10.000          1st Qu.: 9.000          0:242
## Median :10.000          Median :10.000
## Mean    : 9.696          Mean     : 9.403
## 3rd Qu.:10.000          3rd Qu.:10.000
## Max.     :10.000          Max.      :10.000
##
## calculated_host_listings_count calculated_host_listings_count_entire_homes
## Min.     : 1.00          Min.      : 0.000
## 1st Qu.: 1.00          1st Qu.: 0.000
## Median : 2.00          Median : 1.000
## Mean    : 9.53          Mean     : 7.891
## 3rd Qu.: 9.00          3rd Qu.: 7.000
## Max.     :69.00          Max.      :69.000
##
## calculated_host_listings_count_private_rooms reviews_per_month
## Min.     : 0.000          Min.      :0.080
## 1st Qu.: 0.000          1st Qu.:0.530
## Median : 0.000          Median :1.000
## Mean    : 1.624          Mean     :1.648
## 3rd Qu.: 2.000          3rd Qu.:2.260
## Max.     :11.000          Max.      :8.700
##
```

The variables neighbourhood_group_cleansed and bathrooms are useless. They contain NAs only.

```
# Regression with suspicious points
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.0.4
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-1
```

```
#select the best lambda with cross validation
grid = 10^ seq (10,-2, length =100)

mse <- function(true,pred){
  sse <- mean((pred-true)^2)
}

selected_opt <- opt[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_opt<- model.matrix(price~.,selected_opt)[,-1]
y_opt <- selected_opt$price
```



```

selected_train <- train[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_train<- model.matrix(price~.,selected_train)[,-1]
y_train <- selected_train$price

selected_test <- test[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_test<- model.matrix(price~.,selected_test)[,-1]
y_test <- selected_test$price

set.seed (100)
cv_lasso <- cv.glmnet(x_opt, y_opt, alpha = 1, lambda = grid)
optimal_lambda <- cv_lasso$lambda.min
optimal_lambda

```

```
## [1] 0.6579332
```

```

#train the model
lasso_mod =glmnet(x_train,y_train,alpha =1, lambda =optimal_lambda)
summary(lasso_mod)

```

```

##          Length Class      Mode
## a0          1    -none-   numeric
## beta        55    dgCMatrx S4
## df           1    -none-   numeric
## dim          2    -none-   numeric
## lambda       1    -none-   numeric
## dev.ratio    1    -none-   numeric
## nulldev      1    -none-   numeric
## npasses      1    -none-   numeric
## jerr         1    -none-   numeric
## offset       1    -none-   logical
## call         5    -none-   call
## nobs         1    -none-   numeric

```

```

#make predictions on the test set
lasso_pred <- predict(lasso_mod, s = optimal_lambda, newx = x_test)
lasso_coef <- predict(lasso_mod,s=optimal_lambda,type="coefficients")
lasso_coef

```

```

## 56 x 1 sparse Matrix of class "dgCMatrx"
##                                     1
## (Intercept)                        8.792078e+03
## host_response_timewithin a day     -8.692954e+01
## host_response_timewithin a few hours 9.054449e+00
## host_response_timewithin an hour   -6.837487e+01
## host_response_rate                  7.209731e-01
## host_acceptance_rate               2.474053e-01
## host_is_superhost0                 -6.440685e+01
## host_listings_count                -2.795725e-01

```

```

## host_total_listings_count -6.297941e-13
## host_has_profile_pic0 -1.089051e+02
## host_identity_verified0 -1.569526e+01
## neighbourhood_cleansedBarrhaven -1.754701e+01
## neighbourhood_cleansedBay .
## neighbourhood_cleansedBeacon Hill-Cyrville -4.706946e+01
## neighbourhood_cleansedCapital -2.647551e+00
## neighbourhood_cleansedCollege -2.695608e+01
## neighbourhood_cleansedCumberland -1.717832e+01
## neighbourhood_cleansedGloucester-South Nepean -3.819953e+01
## neighbourhood_cleansedGloucester-Southgate -2.588383e+01
## neighbourhood_cleansedInnes -2.825754e+01
## neighbourhood_cleansedKanata North .
## neighbourhood_cleansedKanata South -1.056111e+01
## neighbourhood_cleansedKitchissippi 1.524314e+02
## neighbourhood_cleansedKnoxdale-Merivale -4.074176e+01
## neighbourhood_cleansedOrleans -5.384571e+01
## neighbourhood_cleansedOsgoode -2.110030e+01
## neighbourhood_cleansedRideau-Goulbourn 2.409357e+01
## neighbourhood_cleansedRideau-Rockcliffe -1.565087e+01
## neighbourhood_cleansedRideau-Vanier .
## neighbourhood_cleansedRiver -1.582774e+01
## neighbourhood_cleansedSomerset 9.435319e+01
## neighbourhood_cleansedStittsville-Kanata West .
## neighbourhood_cleansedWest Carleton-March 2.952699e+01
## latitude .
## longitude 1.166092e+02
## accommodates 5.861266e+00
## bedrooms 3.588107e+01
## beds -9.975633e+00
## minimum_nights 1.986073e-01
## maximum_nights -7.260631e-08
## has_availability0 5.249541e+01
## number_of_reviews .
## number_of_reviews_ltm 4.004092e+00
## number_of_reviews_l30d 1.780602e+00
## review_scores_rating .
## review_scores_accuracy .
## review_scores_cleanliness 1.965311e+01
## review_scores_checkin 1.909110e+00
## review_scores_communication 2.567770e+01
## review_scores_location -1.041705e+01
## review_scores_value -2.476359e+01
## instant_bookable0 -1.215086e+01
## calculated_host_listings_count .
## calculated_host_listings_count_entire_homes .
## calculated_host_listings_count_private_rooms 7.304412e-01
## reviews_per_month -2.842106e+01

```

```

testerr <- mse(y_test,lasso_pred)
testerr

```

```
## [1] 7629.632
```

```
#normal regression on dataset
```

```
mod <- lm(price~.,data=selected_train)
```

```
summary(mod)
```

```
##
## Call:
## lm(formula = price ~ ., data = selected_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -374.7  -84.5  -20.9   34.9  9487.0
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value
## (Intercept)    2.267e+04  5.856e+04   0.387
## host_response_timewithin a day    -1.694e+02  2.257e+02  -0.750
## host_response_timewithin a few hours    -7.231e+01  2.397e+02  -0.302
## host_response_timewithin an hour    -1.525e+02  2.326e+02  -0.655
## host_response_rate      1.485e+00  2.258e+00   0.658
## host_acceptance_rate    2.950e-01  7.010e-01   0.421
## host_is_superhost0    -6.698e+01  3.775e+01  -1.775
## host_listings_count    -3.702e-01  1.394e+00  -0.266
## host_total_listings_count              NA              NA              NA
## host_has_profile_pic0    -1.377e+02  4.828e+02  -0.285
## host_identity_verified0    -1.633e+01  4.621e+01  -0.353
## neighbourhood_cleansedBarrhaven    -3.454e+01  1.991e+02  -0.173
## neighbourhood_cleansedBay    -4.819e+00  1.593e+02  -0.030
## neighbourhood_cleansedBeacon Hill-Cyrville    -9.707e+01  1.378e+02  -0.704
## neighbourhood_cleansedCapital    -3.252e+01  1.047e+02  -0.311
## neighbourhood_cleansedCollege    -4.181e+01  1.562e+02  -0.268
## neighbourhood_cleansedCumberland    -9.952e+01  1.888e+02  -0.527
## neighbourhood_cleansedGloucester-South Nepean    -5.978e+01  1.481e+02  -0.404
## neighbourhood_cleansedGloucester-Southgate    -6.779e+01  1.361e+02  -0.498
## neighbourhood_cleansedInnes    -9.613e+01  1.588e+02  -0.605
## neighbourhood_cleansedKanata North      1.541e+01  2.400e+02   0.064
## neighbourhood_cleansedKanata South    -5.951e+00  2.216e+02  -0.027
## neighbourhood_cleansedKitchissippi      1.362e+02  1.137e+02   1.198
## neighbourhood_cleansedKnoxdale-Merivale    -6.796e+01  1.471e+02  -0.462
## neighbourhood_cleansedOrleans    -1.219e+02  1.776e+02  -0.686
## neighbourhood_cleansedOsgoode    -6.827e+01  2.489e+02  -0.274
## neighbourhood_cleansedRideau-Goulbourn      2.166e+01  2.270e+02   0.095
## neighbourhood_cleansedRideau-Rockcliffe    -5.670e+01  1.206e+02  -0.470
## neighbourhood_cleansedRideau-Vanier    -2.723e+01  1.009e+02  -0.270
## neighbourhood_cleansedRiver    -3.960e+01  1.212e+02  -0.327
## neighbourhood_cleansedSomerset      7.493e+01  1.040e+02   0.720
## neighbourhood_cleansedStittsville-Kanata West    3.041e+01  2.516e+02   0.121
## neighbourhood_cleansedWest Carleton-March    8.231e+01  3.025e+02   0.272
## latitude      4.324e+01  8.258e+02   0.052
## longitude     3.254e+02  6.211e+02   0.524
## accommodates    7.076e+00  1.473e+01   0.480
## bedrooms      3.540e+01  3.160e+01   1.120
## beds    -1.144e+01  2.208e+01  -0.518
## minimum_nights    2.013e-01  3.579e-01   0.562
```

## maximum_nights	-8.600e-08	2.278e-07	-0.377
## has_availability0	7.863e+01	1.360e+02	0.578
## number_of_reviews	4.012e-02	5.829e-01	0.069
## number_of_reviews_ltm	4.311e+00	2.292e+00	1.881
## number_of_reviews_l30d	1.718e+00	2.902e+01	0.059
## review_scores_rating	-4.180e-01	5.755e+00	-0.073
## review_scores_accuracy	-1.496e-02	3.852e+01	0.000
## review_scores_cleanliness	2.376e+01	3.550e+01	0.669
## review_scores_checkin	1.542e+00	5.243e+01	0.029
## review_scores_communication	2.811e+01	5.707e+01	0.493
## review_scores_location	-1.091e+01	3.091e+01	-0.353
## review_scores_value	-2.596e+01	3.504e+01	-0.741
## instant_bookable0	-1.869e+01	4.196e+01	-0.445
## calculated_host_listings_count	-8.669e+01	1.204e+02	-0.720
## calculated_host_listings_count_entire_homes	8.674e+01	1.203e+02	0.721
## calculated_host_listings_count_private_rooms	8.994e+01	1.215e+02	0.740
## reviews_per_month	-3.359e+01	3.401e+01	-0.988
##	Pr(> t)		
## (Intercept)	0.6987		
## host_response_timewithin a day	0.4532		
## host_response_timewithin a few hours	0.7630		
## host_response_timewithin an hour	0.5123		
## host_response_rate	0.5109		
## host_acceptance_rate	0.6740		
## host_is_superhost0	0.0763		
## host_listings_count	0.7906		
## host_total_listings_count	NA		
## host_has_profile_pic0	0.7756		
## host_identity_verified0	0.7239		
## neighbourhood_cleansedBarrhaven	0.8623		
## neighbourhood_cleansedBay	0.9759		
## neighbourhood_cleansedBeacon Hill-Cyrville	0.4814		
## neighbourhood_cleansedCapital	0.7562		
## neighbourhood_cleansedCollege	0.7890		
## neighbourhood_cleansedCumberland	0.5983		
## neighbourhood_cleansedGloucester-South Nepean	0.6867		
## neighbourhood_cleansedGloucester-Southgate	0.6186		
## neighbourhood_cleansedInnes	0.5451		
## neighbourhood_cleansedKanata North	0.9488		
## neighbourhood_cleansedKanata South	0.9786		
## neighbourhood_cleansedKitchissippi	0.2312		
## neighbourhood_cleansedKnoxdale-Merivale	0.6443		
## neighbourhood_cleansedOrleans	0.4926		
## neighbourhood_cleansedOsgoode	0.7839		
## neighbourhood_cleansedRideau-Goulbourn	0.9240		
## neighbourhood_cleansedRideau-Rockcliffe	0.6383		
## neighbourhood_cleansedRideau-Vanier	0.7874		
## neighbourhood_cleansedRiver	0.7438		
## neighbourhood_cleansedSomerset	0.4716		
## neighbourhood_cleansedStittsville-Kanata West	0.9038		
## neighbourhood_cleansedWest Carleton-March	0.7856		
## latitude	0.9583		
## longitude	0.6004		
## accommodates	0.6310		

```
## bedrooms                                0.2630
## beds                                    0.6045
## minimum_nights                          0.5739
## maximum_nights                         0.7059
## has_availability0                       0.5634
## number_of_reviews                      0.9451
## number_of_reviews_ltm                  0.0604
## number_of_reviews_l30d                 0.9528
## review_scores_rating                   0.9421
## review_scores_accuracy                 0.9997
## review_scores_cleanliness              0.5034
## review_scores_checkin                  0.9765
## review_scores_communication            0.6224
## review_scores_location                 0.7243
## review_scores_value                    0.4590
## instant_bookable0                     0.6562
## calculated_host_listings_count         0.4716
## calculated_host_listings_count_entire_homes 0.4711
## calculated_host_listings_count_private_rooms 0.4595
## reviews_per_month                     0.3236
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 477.6 on 873 degrees of freedom
## Multiple R-squared:  0.04063,    Adjusted R-squared:  -0.01871
## F-statistic: 0.6847 on 54 and 873 DF,  p-value: 0.9597
```

```
reg_pred <- predict(mod,selected_test)
```

```
## Warning in predict.lm(mod, selected_test): prediction from a rank-deficient fit
## may be misleading
```

```
testerr2 <- mse(selected_test$price,reg_pred)
testerr2
```

```
## [1] 8782.083
```

```
#Regression without suspicious points
#Removing suspicious values
train_new <- train

i1 <- which(train_new$host_listings_count > 200)
train_new <- train_new[-i1,]

i2 <- which(train_new$price > 9000)
train_new <- train_new[-i2,]

i3 <- which(train_new$beds > 8)
train_new <- train_new[-i3,]

i4 <- which(train_new$number_of_reviews > 500)
train_new <- train_new[-i4,]
```

```

i5 <- which(train_new$maximum_nights> 3000)
train_new <- train_new[-i5,]

i6 <- which(train_new$number_of_reviews_ltm > 100)
train_new <- train_new[-i6,]

i7 <- which(train_new$number_of_reviews_l30d > 4)
train_new <- train_new[-i7,]

i8 <- which(train_new$room_type=="Hotel room")
train_new <- train_new[-i8,]
train_new$room_type <- factor(train_new$room_type)

#normal regression

selected_train2 <- train_new[,c(-1,-2,-7,-15,-20,-28,-29)]
selected_test2 <- test[,c(-1,-2,-7,-15,-20,-28,-29)]
mod2 <- lm(price~.,data=selected_train2)
reg_pred2 <- predict(mod2,selected_test2)

## Warning in predict.lm(mod2, selected_test2): prediction from a rank-deficient
## fit may be misleading

testerr3 <- mse(selected_test2$price,reg_pred2)
testerr3

## [1] 9412.505

#lasso regression
x_train2<- model.matrix(price~.,selected_train2)[,-1]
y_train2 <- selected_train2$price

x_test2<- model.matrix(price~.,selected_test2)[,-1]
y_test2 <- selected_test2$price

#make predictions on the test set
lasso_mod2 =glmnet(x_train2,y_train2,alpha =1, lambda =optimal_lambda)
lasso_pred2 <- predict(lasso_mod2, s = optimal_lambda, newx = x_test2)
lasso_coef2 <- predict(lasso_mod2,s=optimal_lambda,type="coefficients")
lasso_coef2

## 58 x 1 sparse Matrix of class "dgCMatrix"
##
## (Intercept) 362.59935968
## host_response_timewithin a day -24.28236218
## host_response_timewithin a few hours 69.69136419
## host_response_timewithin an hour -22.36362625
## host_response_rate 0.39479992
## host_acceptance_rate 0.16473544
## host_is_superhost0 -46.10799535
## host_listings_count .

```

```

## host_total_listings_count .
## host_has_profile_pic0 .
## host_identity_verified0 -6.78262309
## neighbourhood_cleansedBarrhaven -5.81670075
## neighbourhood_cleansedBay -8.88591034
## neighbourhood_cleansedBeacon Hill-Cyrville -42.06722246
## neighbourhood_cleansedCapital .
## neighbourhood_cleansedCollege -34.53889260
## neighbourhood_cleansedCumberland -10.36678978
## neighbourhood_cleansedGloucester-South Nepean -3.18289886
## neighbourhood_cleansedGloucester-Southgate -8.73129313
## neighbourhood_cleansedInnes -32.52841798
## neighbourhood_cleansedKanata North -14.23714023
## neighbourhood_cleansedKanata South -6.25143930
## neighbourhood_cleansedKitchissippi 127.52626268
## neighbourhood_cleansedKnoxdale-Merivale -63.82631203
## neighbourhood_cleansedOrleans -15.71073765
## neighbourhood_cleansedOsgoode .
## neighbourhood_cleansedRideau-Goulbourn 13.05757041
## neighbourhood_cleansedRideau-Rockcliffe .
## neighbourhood_cleansedRideau-Vanier -15.60361123
## neighbourhood_cleansedRiver -13.62803857
## neighbourhood_cleansedSomerset 8.16855131
## neighbourhood_cleansedStittsville-Kanata West .
## neighbourhood_cleansedWest Carleton-March 29.69534078
## latitude 62.31218362
## longitude 42.05976715
## room_typePrivate room -92.98318120
## room_typeShared room -130.96938140
## accommodates 4.25517822
## bedrooms 26.80813998
## beds -4.42445324
## minimum_nights 0.11177620
## maximum_nights 0.02352460
## has_availability0 47.62759260
## number_of_reviews -0.06133796
## number_of_reviews_ltm -0.05625324
## number_of_reviews_l30d .
## review_scores_rating -0.75816715
## review_scores_accuracy 7.67340012
## review_scores_cleanliness 17.03538570
## review_scores_checkin 6.89244771
## review_scores_communication 19.66233872
## review_scores_location -6.09738229
## review_scores_value -27.19507325
## instant_bookable0 -33.99588358
## calculated_host_listings_count .
## calculated_host_listings_count_entire_homes -0.70819760
## calculated_host_listings_count_private_rooms 15.86610128
## reviews_per_month -17.43142192

```

```

testerr4 <- mse(y_test2,lasso_pred2)
testerr4

```

```
## [1] 8216.946
```

```
testerr;testerr2;testerr3;testerr4
```

```
## [1] 7629.632
```

```
## [1] 8782.083
```

```
## [1] 9412.505
```

```
## [1] 8216.946
```

Without removing susicious points:

The test error for standard regression is 8782.0830349.

The test error for Lasso regression is 7629.6322779.

Removed suspicious points:

The test error for standad regression is 9412.504995.

The test error for Lasso regression is 8216.9460501.

In both cases, the test error for Lasso Regression is smaller than the test error for standard regression. But the test errors for the models containing the suspicious points are smaller than the corresponding test errors for the models in which the suspicious points have been removed.

It appears that lasso regression for the model containing the suspicious points is the best model.

It appears that removing the data points caused the sample size to decrease and thus might have inadvertently led to an increase in the test error rate.

```
text <- rep("",928)
for(i in 1:928){
  text[i] <- paste(train$description[i],train$neighborhood_overview[i],train$amenities[i],sep="")
}
n = length(text)
n
```

```
## [1] 928
```

```
w1 <- grepl("bedroom", text, ignore.case = TRUE)
table(w1)/n
```

```
## w1
##      FALSE      TRUE
## 0.3415948 0.6584052
```

```
w2 <- grepl("Wifi", text, ignore.case = TRUE)
table(w2)/n
```



```
## w2
##      FALSE      TRUE
## 0.01831897 0.98168103
```

```
w3<- grepl("heating", text, ignore.case = TRUE)
table(w3)/n
```

```
## w3
##      FALSE      TRUE
## 0.0237069 0.9762931
```

```
w4 <- grepl("walk", text, ignore.case = TRUE)
table(w4)/n
```

```
## w4
##      FALSE      TRUE
## 0.3340517 0.6659483
```

```
w5 <- grepl("university", text, ignore.case = TRUE)
table(w5)/n
```

```
## w5
##      FALSE      TRUE
## 0.8706897 0.1293103
```

```
w6 <- grepl("Refrigerator", text, ignore.case = TRUE)
table(w6)/n
```

```
## w6
##      FALSE      TRUE
## 0.3038793 0.6961207
```

```
w7 <- grepl("downtown", text, ignore.case = TRUE)
table(w7)/n
```

```
## w7
##      FALSE      TRUE
## 0.4353448 0.5646552
```

```
w8 <- grepl("parking", text, ignore.case = TRUE)
table(w8)/n
```

```
## w8
##  FALSE   TRUE
## 0.0625 0.9375
```

```
w9 <- grepl("water", text, ignore.case = TRUE)
table(w9)/n
```

```
## w9
##      FALSE      TRUE
## 0.1831897 0.8168103
```

```
w10 <- grepl("alarm", text, ignore.case = TRUE)
table(w10)/n
```

```
## w10
##      FALSE      TRUE
## 0.0237069 0.9762931
```

```
w11 <- grepl("bus",text,ignore.case = TRUE)
table(w11)/n
```

```
## w11
##      FALSE      TRUE
## 0.5344828 0.4655172
```

```
w12 <- grepl("quiet",text,ignore.case = TRUE)
table(w11)/n
```

```
## w11
##      FALSE      TRUE
## 0.5344828 0.4655172
```

```
bedroom <- ifelse(w1,1,0)
wifi <- ifelse(w2,1,0)
heating <- ifelse(w3,1,0)
walk <- ifelse(w4,1,0)
university <- ifelse(w5,1,0)
refrigerator <- ifelse(w6,1,0)
downtown <- ifelse(w7,1,0)
parking <- ifelse(w8,1,0)
water <- ifelse(w9,1,0)
alarm <- ifelse(w10,1,0)
bus <- ifelse(w11,1,0)
quiet <- ifelse(w12,1,0)

train$bedroom <- bedroom
train$wifi <- wifi
train$heating <- heating
train$walk <- walk
train$university <- university
train$refrigerator <- refrigerator
train$downtown <- downtown
train$parking <- parking
train$water <- water
train$alarm <- alarm
train$bus <- bus
train$quiet <- quiet
```

We choose the above 12 words.
Each $p(w)$ is obtained in the tables (in the column TRUE).

```

#modify the test dataset
texttest <- rep("",404)
for(i in 1:404){
  texttest[i] <- paste(test$description[i],test$neighborhood_overview[i],test$amenities[i],sep="")
}

w11 <- grepl("bedroom", texttest, ignore.case = TRUE)
w22 <- grepl("Wifi", texttest, ignore.case = TRUE)
w33<- grepl("heating", texttest, ignore.case = TRUE)
w44 <- grepl("walk", texttest, ignore.case = TRUE)
w55 <- grepl("university", texttest, ignore.case = TRUE)
w66 <- grepl("Refrigerator", texttest, ignore.case = TRUE)
w77 <- grepl("downtown", texttest, ignore.case = TRUE)
w88 <- grepl("parking", texttest, ignore.case = TRUE)
w99 <- grepl("water", texttest, ignore.case = TRUE)
w100 <- grepl("alarm", texttest, ignore.case = TRUE)
w111 <- grepl("bus",texttest,ignore.case = TRUE)
w122 <- grepl("quiet",texttest,ignore.case = TRUE)

bedroom <- ifelse(w11,1,0)
wifi <- ifelse(w22,1,0)
heating <- ifelse(w33,1,0)
walk <- ifelse(w44,1,0)
university <- ifelse(w55,1,0)
refrigerator <- ifelse(w66,1,0)
downtown <- ifelse(w77,1,0)
parking <- ifelse(w88,1,0)
water <- ifelse(w99,1,0)
alarm <- ifelse(w100,1,0)
bus <- ifelse(w111,1,0)
quiet <- ifelse(w122,1,0)

test$bedroom <- bedroom
test$wifi <- wifi
test$heating <- heating
test$walk <- walk
test$university <- university
test$refrigerator <- refrigerator
test$downtown <- downtown
test$parking <- parking
test$water <- water
test$alarm <- alarm
test$bus <- bus
test$quiet <- quiet
#modify the optimization dataset

textopt <- rep("",848)
for(i in 1:848){
  textopt[i] <- paste(opt$description[i],opt$neighborhood_overview[i],opt$amenities[i],sep="")
}

o1 <- grepl("bedroom", textopt, ignore.case = TRUE)
o2 <- grepl("Wifi", textopt, ignore.case = TRUE)

```

```

o3<- grepl("heating", textopt, ignore.case = TRUE)
o4 <- grepl("walk", textopt, ignore.case = TRUE)
o5 <- grepl("university", textopt, ignore.case = TRUE)
o6 <- grepl("Refrigerator", textopt, ignore.case = TRUE)
o7 <- grepl("downtown", textopt, ignore.case = TRUE)
o8 <- grepl("parking", textopt, ignore.case = TRUE)
o9 <- grepl("water", textopt, ignore.case = TRUE)
o10 <- grepl("alarm", textopt, ignore.case = TRUE)
o11 <- grepl("bus",textopt,ignore.case = TRUE)
o12 <- grepl("quiet",textopt,ignore.case = TRUE)

bedroom <- ifelse(o1,1,0)
wifi <- ifelse(o2,1,0)
heating <- ifelse(o3,1,0)
walk <- ifelse(o4,1,0)
university <- ifelse(o5,1,0)
refrigerator <- ifelse(o6,1,0)
downtown <- ifelse(o7,1,0)
parking <- ifelse(o8,1,0)
water <- ifelse(o9,1,0)
alarm <- ifelse(o10,1,0)
bus <- ifelse(o11,1,0)
quiet <- ifelse(o12,1,0)

opt$bedroom <- bedroom
opt$wifi <- wifi
opt$heating <- heating
opt$walk <- walk
opt$university <- university
opt$refrigerator <- refrigerator
opt$downtown <- downtown
opt$parking <- parking
opt$water <- water
opt$alarm <- alarm
opt$bus <- bus
opt$quiet <- quiet

#standard regression model
selected_train <- train[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_train<- model.matrix(price~.,selected_train)[-1]
y_train <- selected_train$price

selected_test <- test[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_test<- model.matrix(price~.,selected_test)[-1]
y_test <- selected_test$price

selected_opt <- opt[, c(-1,-2,-7,-15,-16,-20,-28,-29)]
x_opt<- model.matrix(price~.,selected_opt)[-1]
y_opt <- selected_opt$price

#Find the new lambda
set.seed (100)

```

```
las <- cv.glmnet(x_opt, y_opt, alpha = 1, lambda = grid)
optimal <- las$lambda.min
optimal
```

```
## [1] 0.869749
```

```
#train the model
lassomod =glmnet(x_train,y_train,alpha=1, lambda =optimal)
summary(lassomod)
```

```
##          Length Class      Mode
## a0          1    -none-   numeric
## beta        67    dgCMatrx S4
## df           1    -none-   numeric
## dim          2    -none-   numeric
## lambda       1    -none-   numeric
## dev.ratio    1    -none-   numeric
## nulldev      1    -none-   numeric
## npasses      1    -none-   numeric
## jerr         1    -none-   numeric
## offset       1    -none-   logical
## call         5    -none-   call
## nobs         1    -none-   numeric
```

```
#make predictions on the test set
lasso_pred <- predict(lassomod, s = optimal, newx = x_test)
lasso_coef <- predict(lassomod,s=optimal,type="coefficients")
lasso_coef
```

```
## 68 x 1 sparse Matrix of class "dgCMatrx"
##
## (Intercept) 8.122593e+03
## host_response_timewithin a day -5.312847e+01
## host_response_timewithin a few hours 4.266451e+01
## host_response_timewithin an hour -3.336827e+01
## host_response_rate 4.430360e-01
## host_acceptance_rate 2.700334e-01
## host_is_superhost0 -6.671329e+01
## host_listings_count -2.892500e-01
## host_total_listings_count -4.284730e-13
## host_has_profile_pic0 -1.110841e+02
## host_identity_verified0 -1.380409e+01
## neighbourhood_cleansedBarrhaven -7.329558e+00
## neighbourhood_cleansedBay .
## neighbourhood_cleansedBeacon Hill-Cyrville -4.234317e+01
## neighbourhood_cleansedCapital -3.307401e-01
## neighbourhood_cleansedCollege -3.077579e+01
## neighbourhood_cleansedCumberland -9.701915e+00
## neighbourhood_cleansedGloucester-South Nepean -3.457931e+01
## neighbourhood_cleansedGloucester-Southgate -2.363459e+01
## neighbourhood_cleansedInnes -2.588075e+01
## neighbourhood_cleansedKanata North .
```

```

## neighbourhood_cleansedKanata South -1.547461e+01
## neighbourhood_cleansedKitchissippi 1.483358e+02
## neighbourhood_cleansedKnoxdale-Merivale -3.285513e+01
## neighbourhood_cleansedOrleans -5.783579e+01
## neighbourhood_cleansedOsgoode -2.186176e+01
## neighbourhood_cleansedRideau-Goulbourn 1.970330e+01
## neighbourhood_cleansedRideau-Rockcliffe -8.638366e+00
## neighbourhood_cleansedRideau-Vanier 2.747876e+00
## neighbourhood_cleansedRiver -1.926635e+00
## neighbourhood_cleansedSomerset 9.211956e+01
## neighbourhood_cleansedStittsville-Kanata West .
## neighbourhood_cleansedWest Carleton-March 2.513699e+01
## latitude .
## longitude 1.078330e+02
## accommodates 6.188338e+00
## bedrooms 3.434256e+01
## beds -9.768057e+00
## minimum_nights 2.250069e-01
## maximum_nights -7.094830e-08
## has_availability0 5.550456e+01
## number_of_reviews -1.726663e-02
## number_of_reviews_ltm 3.847302e+00
## number_of_reviews_l30d 1.856676e-01
## review_scores_rating .
## review_scores_accuracy .
## review_scores_cleanliness 2.209868e+01
## review_scores_checkin 3.751949e+00
## review_scores_communication 2.300307e+01
## review_scores_location -1.071529e+01
## review_scores_value -2.710528e+01
## instant_bookable0 -1.261257e+01
## calculated_host_listings_count .
## calculated_host_listings_count_entire_homes .
## calculated_host_listings_count_private_rooms 1.408283e+00
## reviews_per_month -2.662334e+01
## bedroom -1.350833e+01
## wifi 5.425367e+00
## heating 2.790146e+01
## walk 2.447594e+01
## university -3.616422e+01
## refrigerator -1.464859e+00
## downtown -1.229204e+01
## parking 1.899674e+01
## water -3.803156e+00
## alarm -2.651582e+01
## bus -2.315646e+01
## quiet .

```

```

test_error2 <- mse(y_test,lasso_pred)
test_error2

```

```
## [1] 8195.195
```

#Standard regression

```
mod3 <- lm(price~.,data=selected_train)
summary(mod3)
```

```
##
## Call:
## lm(formula = price ~ ., data = selected_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -399.8  -83.8  -20.2   35.5  9488.5
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value
## (Intercept)    2.796e+04  6.009e+04   0.465
## host_response_timewithin a day    -1.568e+02  2.302e+02  -0.681
## host_response_timewithin a few hours    -6.108e+01  2.446e+02  -0.250
## host_response_timewithin an hour    -1.402e+02  2.375e+02  -0.590
## host_response_rate      1.433e+00  2.306e+00   0.621
## host_acceptance_rate    3.512e-01  7.267e-01   0.483
## host_is_superhost0    -7.055e+01  3.852e+01  -1.831
## host_listings_count    -3.351e-01  1.413e+00  -0.237
## host_total_listings_count              NA              NA              NA
## host_has_profile_pic0    -1.530e+02  4.872e+02  -0.314
## host_identity_verified0    -1.510e+01  4.697e+01  -0.322
## neighbourhood_cleansedBarrhaven    -3.477e+01  2.018e+02  -0.172
## neighbourhood_cleansedBay    -2.349e+00  1.625e+02  -0.014
## neighbourhood_cleansedBeacon Hill-Cyrville    -8.897e+01  1.405e+02  -0.633
## neighbourhood_cleansedCapital    -2.772e+01  1.072e+02  -0.259
## neighbourhood_cleansedCollege    -4.789e+01  1.592e+02  -0.301
## neighbourhood_cleansedCumberland    -9.072e+01  1.915e+02  -0.474
## neighbourhood_cleansedGloucester-South Nepean    -6.794e+01  1.509e+02  -0.450
## neighbourhood_cleansedGloucester-Southgate    -6.937e+01  1.378e+02  -0.504
## neighbourhood_cleansedInnes    -9.417e+01  1.640e+02  -0.574
## neighbourhood_cleansedKanata North      1.275e+01  2.446e+02   0.052
## neighbourhood_cleansedKanata South    -1.824e+01  2.258e+02  -0.081
## neighbourhood_cleansedKitchissippi      1.364e+02  1.163e+02   1.172
## neighbourhood_cleansedKnoxdale-Merivale    -6.512e+01  1.493e+02  -0.436
## neighbourhood_cleansedOrleans    -1.262e+02  1.812e+02  -0.696
## neighbourhood_cleansedOsgoode    -9.141e+01  2.526e+02  -0.362
## neighbourhood_cleansedRideau-Goulbourn      3.600e+00  2.313e+02   0.016
## neighbourhood_cleansedRideau-Rockcliffe    -4.644e+01  1.230e+02  -0.377
## neighbourhood_cleansedRideau-Vanier    -1.559e+01  1.037e+02  -0.150
## neighbourhood_cleansedRiver    -2.454e+01  1.238e+02  -0.198
## neighbourhood_cleansedSomerset      7.827e+01  1.068e+02   0.733
## neighbourhood_cleansedStittsville-Kanata West    1.435e+01  2.558e+02   0.056
## neighbourhood_cleansedWest Carleton-March      9.118e+01  3.083e+02   0.296
## latitude    -5.158e+01  8.418e+02  -0.061
## longitude     3.387e+02  6.349e+02   0.533
## accommodates     8.442e+00  1.508e+01   0.560
## bedrooms       3.408e+01  3.221e+01   1.058
## beds    -1.247e+01  2.239e+01  -0.557
## minimum_nights     2.359e-01  3.621e-01   0.651
```

## maximum_nights	-8.387e-08	2.384e-07	-0.352
## has_availability0	8.529e+01	1.426e+02	0.598
## number_of_reviews	-9.327e-03	5.921e-01	-0.016
## number_of_reviews_ltm	4.190e+00	2.320e+00	1.806
## number_of_reviews_l30d	2.199e-01	2.943e+01	0.007
## review_scores_rating	-6.519e-01	5.860e+00	-0.111
## review_scores_accuracy	-1.215e+00	3.925e+01	-0.031
## review_scores_cleanliness	2.828e+01	3.608e+01	0.784
## review_scores_checkin	4.591e+00	5.360e+01	0.086
## review_scores_communication	2.661e+01	5.783e+01	0.460
## review_scores_location	-1.131e+01	3.174e+01	-0.356
## review_scores_value	-2.841e+01	3.588e+01	-0.792
## instant_bookable0	-1.940e+01	4.263e+01	-0.455
## calculated_host_listings_count	-7.293e+01	1.234e+02	-0.591
## calculated_host_listings_count_entire_homes	7.272e+01	1.233e+02	0.590
## calculated_host_listings_count_private_rooms	7.804e+01	1.245e+02	0.627
## reviews_per_month	-3.045e+01	3.470e+01	-0.877
## bedroom	-1.747e+01	3.544e+01	-0.493
## wifi	1.954e+01	1.452e+02	0.135
## heating	2.566e+01	1.227e+02	0.209
## walk	2.664e+01	3.827e+01	0.696
## university	-4.113e+01	5.252e+01	-0.783
## refrigerator	-6.324e+00	4.509e+01	-0.140
## downtown	-1.496e+01	3.407e+01	-0.439
## parking	2.102e+01	7.337e+01	0.286
## water	-6.987e+00	5.384e+01	-0.130
## alarm	-3.066e+01	1.167e+02	-0.263
## bus	-2.199e+01	3.381e+01	-0.650
## quiet	-3.873e-01	3.464e+01	-0.011
##	Pr(> t)		
## (Intercept)	0.6418		
## host_response_timewithin a day	0.4959		
## host_response_timewithin a few hours	0.8029		
## host_response_timewithin an hour	0.5551		
## host_response_rate	0.5346		
## host_acceptance_rate	0.6290		
## host_is_superhost0	0.0674		
## host_listings_count	0.8126		
## host_total_listings_count	NA		
## host_has_profile_pic0	0.7536		
## host_identity_verified0	0.7479		
## neighbourhood_cleansedBarrhaven	0.8632		
## neighbourhood_cleansedBay	0.9885		
## neighbourhood_cleansedBeacon Hill-Cyrville	0.5267		
## neighbourhood_cleansedCapital	0.7960		
## neighbourhood_cleansedCollege	0.7637		
## neighbourhood_cleansedCumberland	0.6359		
## neighbourhood_cleansedGloucester-South Nepean	0.6526		
## neighbourhood_cleansedGloucester-Southgate	0.6147		
## neighbourhood_cleansedInnes	0.5660		
## neighbourhood_cleansedKanata North	0.9584		
## neighbourhood_cleansedKanata South	0.9356		
## neighbourhood_cleansedKitchissippi	0.2415		
## neighbourhood_cleansedKnoxdale-Merivale	0.6628		


```
## neighbourhood_cleansedOrleans 0.4864
## neighbourhood_cleansedOsgoode 0.7176
## neighbourhood_cleansedRideau-Goulbourn 0.9876
## neighbourhood_cleansedRideau-Rockcliffe 0.7060
## neighbourhood_cleansedRideau-Vanier 0.8805
## neighbourhood_cleansedRiver 0.8430
## neighbourhood_cleansedSomerset 0.4638
## neighbourhood_cleansedStittsville-Kanata West 0.9553
## neighbourhood_cleansedWest Carleton-March 0.7675
## latitude 0.9512
## longitude 0.5938
## accommodates 0.5758
## bedrooms 0.2904
## beds 0.5777
## minimum_nights 0.5150
## maximum_nights 0.7250
## has_availability0 0.5500
## number_of_reviews 0.9874
## number_of_reviews_ltm 0.0713
## number_of_reviews_l30d 0.9940
## review_scores_rating 0.9114
## review_scores_accuracy 0.9753
## review_scores_cleanliness 0.4334
## review_scores_checkin 0.9318
## review_scores_communication 0.6455
## review_scores_location 0.7216
## review_scores_value 0.4286
## instant_bookable0 0.6492
## calculated_host_listings_count 0.5547
## calculated_host_listings_count_entire_homes 0.5556
## calculated_host_listings_count_private_rooms 0.5311
## reviews_per_month 0.3805
## bedroom 0.6223
## wifi 0.8930
## heating 0.8343
## walk 0.4865
## university 0.4338
## refrigerator 0.8885
## downtown 0.6607
## parking 0.7746
## water 0.8968
## alarm 0.7928
## bus 0.5156
## quiet 0.9911
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 480.3 on 861 degrees of freedom
## Multiple R-squared:  0.04302,    Adjusted R-squared:  -0.03034
## F-statistic: 0.5864 on 66 and 861 DF,  p-value: 0.9964
```

```
reg_pred3 <- predict(mod3,selected_test)
```

```
## Warning in predict.lm(mod3, selected_test): prediction from a rank-deficient fit
```

```
## may be misleading
```

```
test_error <- mse(selected_test$price,reg_pred3)
test_error
```

```
## [1] 9814.827
```

The test error for lasso regression is 8195.1953312. It is higher than the test error for lasso regression with suspicious points.

As expected, the test error for the standard regression (9814.8270392) is higher than that for lasso but it is also higher than the test error for standard regression with suspicious points

The new variables added to the model didn't improve the model.